

```
[151]: #lib for extraction ,manipulation,analysis
import numpy as np
import pandas as pd
# for visualisation
import matplotlib.pyplot as plt
import seaborn as sns
# for stats
import scipy.stats
from scipy.stats import shapiro, chi2, normaltest, kstest, zscore
# for vif
from statsmodels.stats.outliers_influence import variance_inflation_factor
# train test split
from sklearn.model_selection import train_test_split
# Linear regression
from sklearn.linear_model import LinearRegression, Lasso,Ridge
# regression evaluation metrics
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
[152]: # importing dataset
data=pd.read_csv(r"C:\Users\deshm\OneDrive\Desktop\ML\03-LinearRegression\mediclaim.csv")
data
```

| | age | sex | bmi | children | smoker | region | charges |
|------|-----|--------|--------|----------|--------|-----------|-------------|
| 0 | 19 | female | 27.900 | 0 | yes | southwest | 16884.92400 |
| 1 | 18 | male | 33.770 | 1 | no | southeast | 1725.55230 |
| 2 | 28 | male | 33.000 | 3 | no | southeast | 4449.46200 |
| 3 | 33 | male | 22.705 | 0 | no | northwest | 21984.47061 |
| 4 | 32 | male | 28.880 | 0 | no | northwest | 3866.85520 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1333 | 50 | male | 30.970 | 3 | no | northwest | 10600.54830 |
| 1334 | 18 | female | 31.920 | 0 | no | northeast | 2205.98080 |
| 1335 | 18 | female | 36.850 | 0 | no | southeast | 1629.83350 |
| 1336 | 21 | female | 25.800 | 0 | no | southwest | 2007.94500 |
| 1337 | 61 | female | 29.070 | 0 | yes | northwest | 29141.36030 |

1338 rows × 7 columns

```
[153]: #EDA for each data analysis
def eda(data):
    print("Shape:",data.shape)
    print("*****")
    print("Size:",data.size)
    print("*****")
    print("INFO:",data.info)
    print("*****")
    print("Describe:",data.describe())
    print("*****")
    print("Dtype:",data.dtypes)
    print("*****")
    print("Checking Null Values:",data.isnull().sum())

eda(data)
```

Shape: (1338, 7)

```
-----
Size: 9366
-----
INFO: <bound method DataFrame.info of   age      sex      bmi  children  smoker     region  charges
0   19  female  27.900      0    yes  southwest  16884.92400
1   18    male  33.770      1     no  southeast  1725.55230
2   28    male  33.000      3     no  southeast  4449.46200
3   33    male  22.705      0     no  northwest  21984.47061
4   32    male  28.880      0     no  northwest  3866.85520
...   ...
1333  50    male  30.970      3     no  northwest  10600.54830
1334  18  female  31.920      0     no  northeast  2205.98080
1335  18  female  36.850      0     no  southeast  1629.83350
1336  21  female  25.800      0     no  southwest  2007.94500
1337  61  female  29.070      0    yes  northwest  29141.36030
```

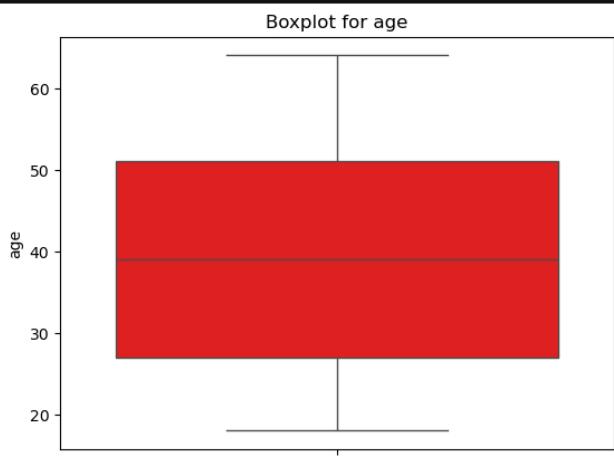
[1338 rows × 7 columns]

```
-----
Describe:           age        bmi  children  charges
count  1338.000000  1338.000000  1338.000000  1338.000000
mean   39.207025  30.663979  1.094918  13270.422265
std    14.049960  6.098187  1.205493  12110.011237
min    18.000000  15.960000  0.000000  1121.873900
25%    27.000000  26.296250  0.000000  4740.287150
50%    39.000000  30.400000  1.000000  9382.033000
75%    51.000000  34.693750  2.000000  16639.912515
max    64.000000  53.130000  5.000000  63770.428010
-----
Dtype: age      int64
       sex      object
       bmi     float64
```



```
#Handling of Outliers :
data.loc[data[col] < LowerTail, col] = LowerTail # all outliers less than lowertail, assigned by lowertail value
data.loc[data[col] > UpperTail, col] = UpperTail # all outliers greater than uppertail, assigned by uppertail value
print("After handling of Outliers data:\n")
print(data.head())
```

[160]: Checking_and_Handling_Of_Outliers(data, "age")



25% Quantile q1 = 27.0
75% Quantile q3 = 51.0
IQR = 24.0

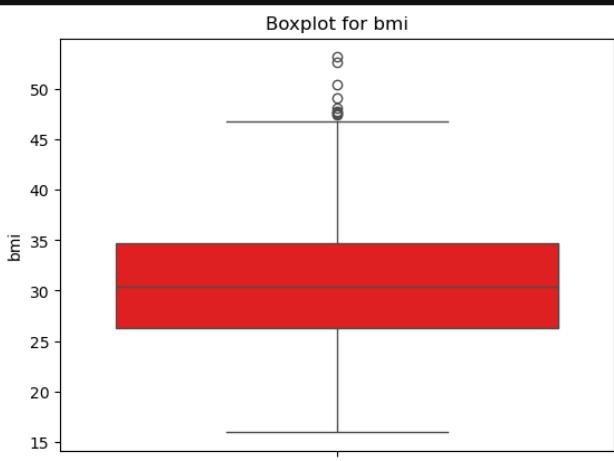
Lower Tail = -9.0
Upper Tail = 87.0

Outliers :
Empty DataFrame
Columns: [age, sex, bmi, children, smoker, region, charges]
Index: []

After handling of Outliers data:

| age | sex | bmi | children | smoker | region | charges |
|-----|-----|--------|----------|--------|--------|-------------|
| 0 | 19 | female | 27.900 | 0 | yes | 16884.92400 |
| 1 | 18 | male | 33.770 | 1 | no | 1725.55230 |
| 2 | 28 | male | 33.000 | 3 | no | 4449.46200 |
| 3 | 33 | male | 22.705 | 0 | no | 21984.47061 |
| 4 | 32 | male | 28.880 | 0 | no | 3866.85520 |

[161]: Checking_and_Handling_Of_Outliers(data, "bmi")



25% Quantile q1 = 26.29625
75% Quantile q3 = 34.69375
IQR = 8.3975

Lower Tail = 13.7
Upper Tail = 47.290000000000006

Outliers :

| age | sex | bmi | children | smoker | region | charges |
|------|-----|--------|----------|--------|--------|-------------|
| 116 | 58 | male | 49.06 | 0 | no | 11381.32540 |
| 286 | 46 | female | 48.07 | 2 | no | 9432.92530 |
| 401 | 47 | male | 47.52 | 1 | no | 8083.91980 |
| 543 | 54 | female | 47.41 | 0 | yes | 63770.42801 |
| 847 | 23 | male | 50.38 | 1 | no | 2438.05520 |
| 860 | 37 | female | 47.60 | 2 | yes | 46113.51100 |
| 1047 | 22 | male | 52.58 | 1 | yes | 44501.39820 |
| 1088 | 52 | male | 47.74 | 1 | no | 9748.91060 |
| 1317 | 18 | male | 53.13 | 0 | no | 1163.46270 |

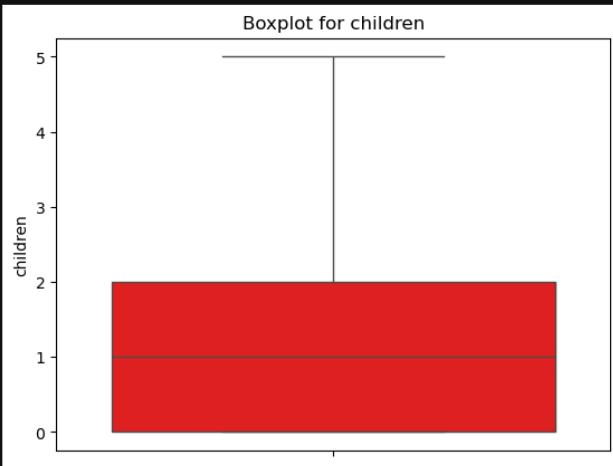
After handling of Outliers data:

```

      age   sex   bmi children smoker    region    charges
0  19  female 27.900      0   yes southwest 16884.92400
1  18    male 33.770      1   no southeast 1725.55230
2  28    male 33.000      3   no southeast 4449.46200
3  33    male 22.705      0   no northwest 21984.47061
4  32    male 28.880      0   no northwest 3866.85520

```

[162]: Checking_and_Handling_Of_Outliers(data, "children")



Lower Tail = -3.0

Upper Tail = 5.0

Outliers :

Empty DataFrame

Columns: [age, sex, bmi, children, smoker, region, charges]

Index: []

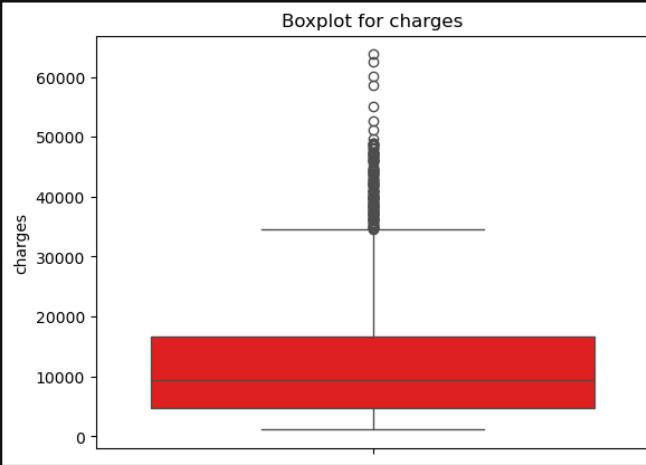
After handling of Outliers data:

```

      age   sex   bmi children smoker    region    charges
0  19  female 27.900      0   yes southwest 16884.92400
1  18    male 33.770      1   no southeast 1725.55230
2  28    male 33.000      3   no southeast 4449.46200
3  33    male 22.705      0   no northwest 21984.47061
4  32    male 28.880      0   no northwest 3866.85520

```

[163]: Checking_and_Handling_Of_Outliers(data, "charges")



Lower Tail = -13109.1508975

Upper Tail = 34489.350562499996

Outliers :

```

      age   sex   bmi children smoker    region    charges
14  27    male 42.130      0   yes southeast 39611.75770
19  30    male 35.300      0   yes southwest 36837.46700
23  34  female 31.920      1   yes northeast 37701.87680
29  31    male 36.300      2   yes southwest 38711.00000
30  22    male 35.600      0   yes southwest 35585.57600
...  ...
1300 45    male 30.360      0   yes southeast 62592.87309
1301 62    male 30.875      3   yes northwest 46718.16325
1303 43    male 27.800      0   yes southwest 37829.72420
1313 19  female 34.700      2   yes southwest 36397.57600
1323 42  female 40.370      2   yes southeast 43896.37630

```

[170]: Checking_and_Handling_Of_Outliers(data, "bmi")

```
[159 rows x 7 columns]
```

```
-----
```

```
After handling of Outliers data:
```

```
age    sex    bmi  children smoker   region   charges
0    19  female  27.900      0    yes southwest 16884.92400
1    18    male  33.770      1    no  southeast 1725.55230
2    28    male  33.000      3    no  southeast 4449.46200
3    33    male  22.705      0    no northwest 21984.47061
4    32    male  28.880      0    no northwest 3866.85520
```

```
[164]: # Analysis of Categorical Columns(Variable)
def Cat_col(data, col):
    unique_values = data[col].unique() # Fixed typo: renamed to unique_values
    value_counts = data[col].value_counts()
    mode = data[col].mode()[0] # Fixed mode access by adding parentheses

    # Enhanced string formatting for clarity
    print(f"Unique Values in '{col}':\n{unique_values}\n")
    print(f"Value Counts in '{col}':\n{value_counts}\n")
    print(f"Mode of '{col}': {mode}\n")

    data[col].value_counts().plot.pie(autopct="%1.1f%%")
    plt.title(f"data-{col} (pie chart)")
    plt.show
```

```
[165]: data.columns
```

```
[165]: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dtype='object')
```

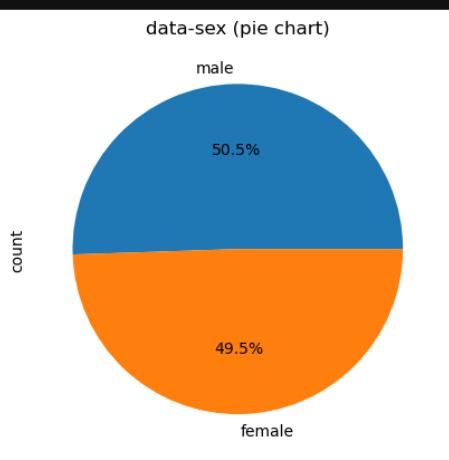
```
[166]: Cat_col(data, "sex")
```

```
Unique Values in 'sex':
['female' 'male']
```

```
Value Counts in 'sex':
```

```
sex
male    676
female   662
Name: count, dtype: int64
```

```
Mode of 'sex': male
```



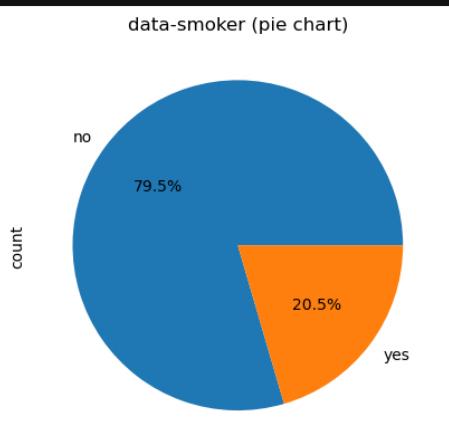
```
[167]: Cat_col(data, "smoker")
```

```
Unique Values in 'smoker':
['yes' 'no']
```

```
Value Counts in 'smoker':
```

```
smoker
no     1064
yes    274
Name: count, dtype: int64
```

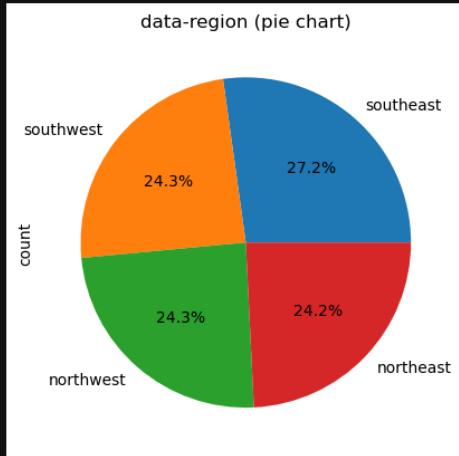
```
Mode of 'smoker': no
```



```
[168]: Cat_col(data, "region")
Unique Values in 'region':
['southwest' 'southeast' 'northwest' 'northeast']

Value Counts in 'region':
region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64

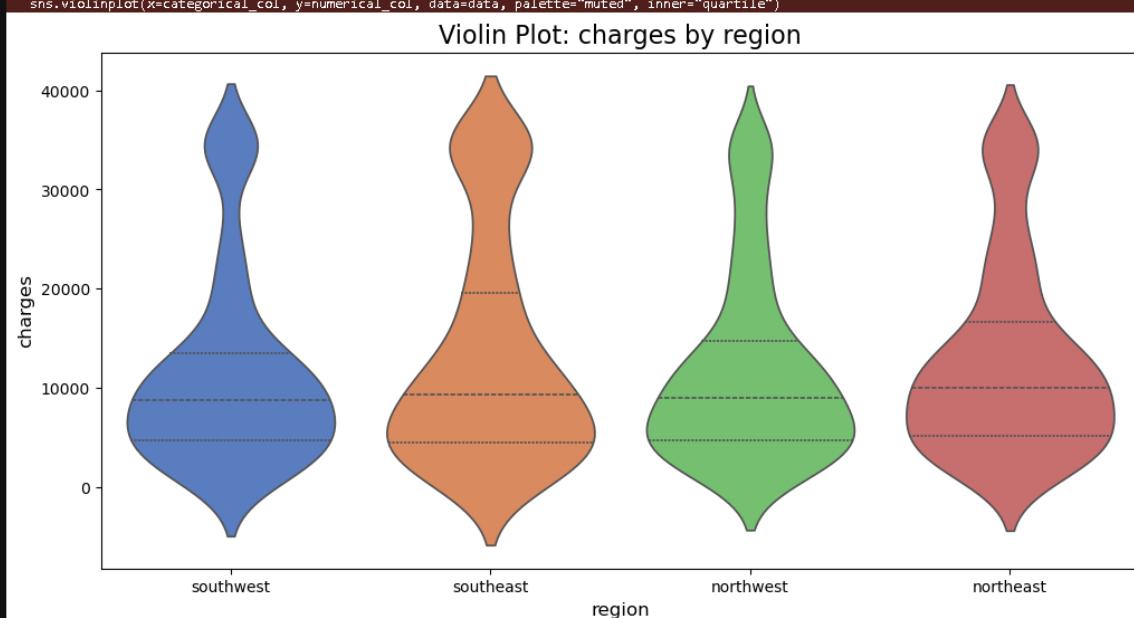
Mode of 'region': southeast
```



```
[169]: # Bivariate analysis of columns
def bivariate_violin_plot(data, categorical_col, numerical_col):
    plt.figure(figsize=(12, 6))
    sns.violinplot(x=categorical_col, y=numerical_col, data=data, palette="muted", inner="quartile")
    plt.title(f'Violin Plot: {numerical_col} by {categorical_col}', fontsize=16)
    plt.xlabel(categorical_col, fontsize=12)
    plt.ylabel(numerical_col, fontsize=12)
    plt.show()
```

```
[170]: bivariate_violin_plot(data, "region", "charges")
C:\Users\deshm\AppData\Local\Temp\ipykernel_2704\3790340880.py:4: FutureWarning:
Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'x' variable to 'hue' and set 'legend=False' for the same effect.

sns.violinplot(x=categorical_col, y=numerical_col, data=data, palette="muted", inner="quartile")
```



```
[171]: # Encoding
data["sex"].unique()

[171]: array(['female', 'male'], dtype=object)

[172]: data["sex_encoded"] = data["sex"].replace({'female':0, 'male':1})
```

```
C:\Users\deshm\AppData\Local\Temp\ipykernel_2704\3863865108.py:1: FutureWarning: Downcasting behavior in 'replace' is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
```

```

[173]: data["sex_encoded"] = data["sex"].replace({'female':0, 'male':1})
[173]: data["smoker"].unique()
[173]: array(['yes', 'no'], dtype=object)

[174]: data["smoker_encoded"] = data["smoker"].replace({'yes':0, 'no':1})

C:\Users\deshm\AppData\Local\Temp\ipykernel_2704\3939261919.py:1: FutureWarning: Downcasting behavior in 'replace' is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
    data["smoker_encoded"] = data["smoker"].replace({'yes':0, 'no':1})

[175]: data["region"].unique()
[175]: array(['southwest', 'southeast', 'northwest', 'northeast'], dtype=object)

[176]: data["region_encoded"] = data["region"].replace({'southwest':0, 'southeast':1, 'northwest':2, 'northeast':3})

C:\Users\deshm\AppData\Local\Temp\ipykernel_2704\1909166453.py:1: FutureWarning: Downcasting behavior in 'replace' is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
    data["region_encoded"] = data["region"].replace({'southwest':0, 'southeast':1, 'northwest':2, 'northeast':3})

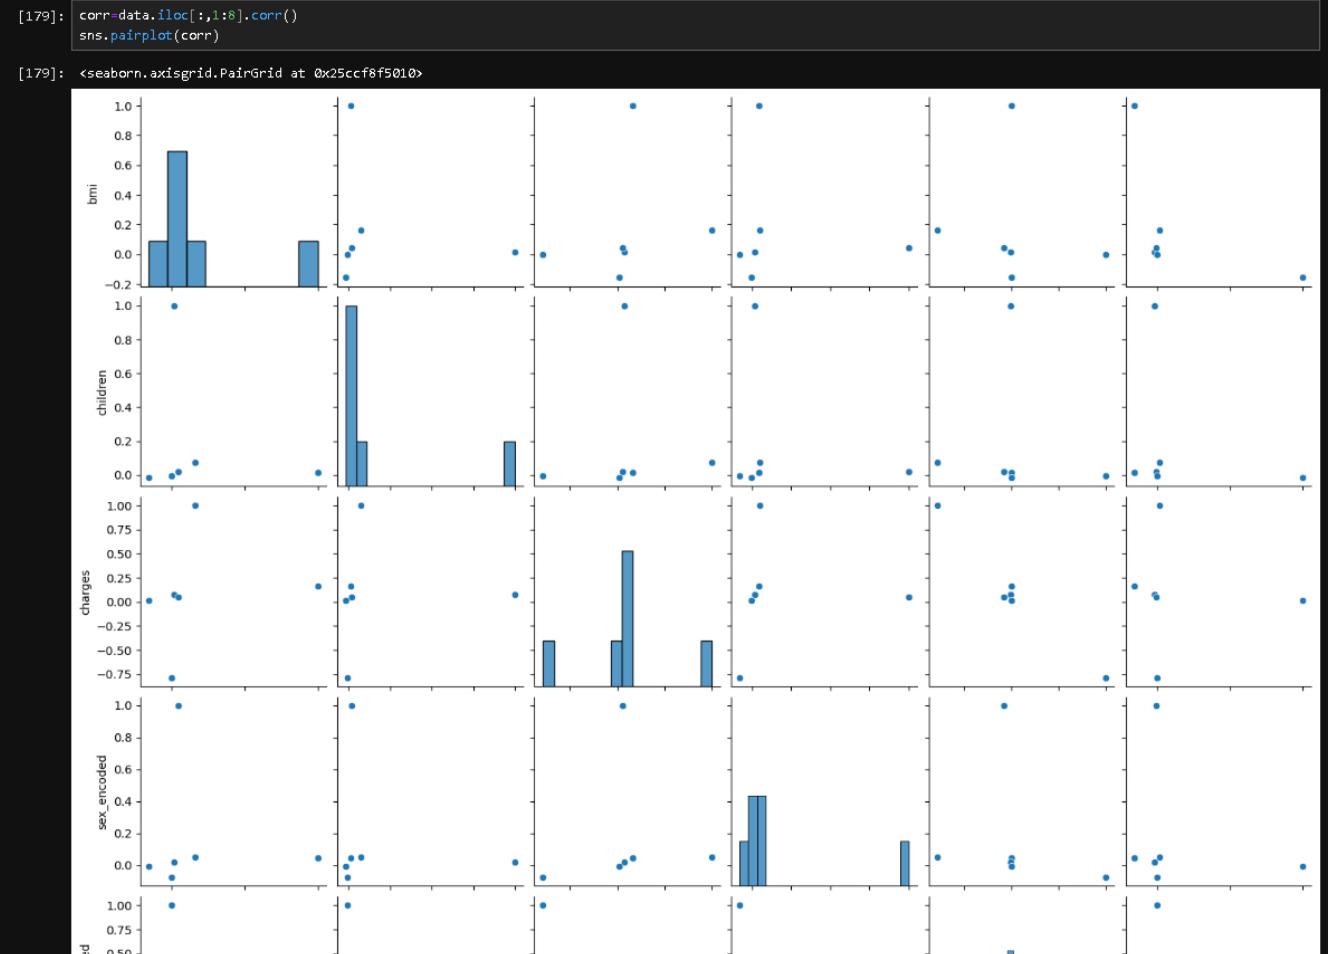
[177]: data.drop(columns=["sex"], inplace=True)
[177]: data.drop(columns=["smoker"], inplace=True)
[177]: data.drop(columns=["region"], inplace=True)

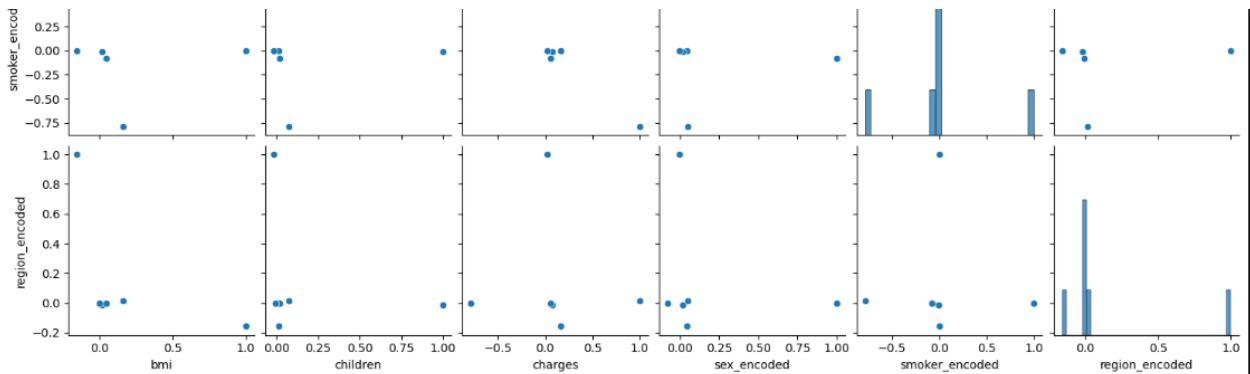
[178]: data

```

| | age | bmi | children | charges | sex_encoded | smoker_encoded | region_encoded |
|------|-----|--------|----------|-------------|-------------|----------------|----------------|
| 0 | 19 | 27.900 | 0 | 16884.92400 | 0 | 0 | 0 |
| 1 | 18 | 33.770 | 1 | 1725.55230 | 1 | 1 | 1 |
| 2 | 28 | 33.000 | 3 | 4449.46200 | 1 | 1 | 1 |
| 3 | 33 | 22.705 | 0 | 21984.47061 | 1 | 1 | 2 |
| 4 | 32 | 28.880 | 0 | 3866.85520 | 1 | 1 | 2 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1333 | 50 | 30.970 | 3 | 10600.54830 | 1 | 1 | 2 |
| 1334 | 18 | 31.920 | 0 | 2205.98080 | 0 | 1 | 3 |
| 1335 | 18 | 36.850 | 0 | 1629.83350 | 0 | 1 | 1 |
| 1336 | 21 | 25.800 | 0 | 2007.94500 | 0 | 1 | 0 |
| 1337 | 61 | 29.070 | 0 | 29141.36030 | 0 | 0 | 2 |

1338 rows × 7 columns





```
[180]: # Train-test split

X = data.drop(columns=['charges']) # Replace 'target_column' with your target variable
Y = data['charges']

xtrain, xtest, ytrain, ytest = train_test_split(X, Y, test_size=0.25, random_state=42)

print(f'Shape of X_train: {xtrain.shape}')
print(f'Shape of X_test: {xtest.shape}')
print(f'Shape of y_train: {ytrain.shape}')
print(f'Shape of y_test: {ytest.shape}')

Shape of X_train: (1003, 6)
Shape of X_test: (335, 6)
Shape of y_train: (1003,)
Shape of y_test: (335,)
```

ML Model Training

```
[182]: lin_reg = LinearRegression()
lin_reg

[182]: +-- LinearRegression ⓘ ??
LinearRegression()

[183]: lin_reg_model = lin_reg.fit(xtrain,ytrain) # Linear regression working, bfl, cost, gradient descent
lin_reg_model

[183]: +-- LinearRegression ⓘ ??
LinearRegression()

[184]: ytrain_pred = lin_reg_model.predict(xtrain)

[184]: array([ 3016.17380033,  5180.34750712, 12707.96514625, ...,
       11448.39181563, 31507.89919888, 11000.5053457 ])
```

```
[185]: # Ensure consistent lengths
print("----- Evaluation for Training Data -----")
print("Shape of ytrain:", ytrain.shape)
print("Shape of ytrain_pred:", ytrain_pred.shape)

# Recompute predictions if necessary
if len(ytrain) != len(ytrain_pred):
    ytrain_pred = ytrain_pred[:len(ytrain)]

# Calculate MSE
mse = mean_squared_error(ytrain, ytrain_pred)
print("Mean Squared Error (MSE):", mse)
print("-" * 80)

mae = mean_absolute_error(ytrain, ytrain_pred)
print("mean absolute error (MAE) : ", mae)
print("-" * 80)

rmse = np.sqrt(mse)
print("Root mean squared error (RMSE) : ", rmse)
print("-" * 80)

r2score = r2_score(ytrain, ytrain_pred)
print("R2 Score : ", r2score)
print("-" * 80)

----- Evaluation for Training Data -----
Shape of ytrain: (1003,)
Shape of ytrain_pred: (1003,)
Mean Squared Error (MSE): 26162370.951339696
-----
mean absolute error (MAE) :  3346.360391711921
-----
Root mean squared error (RMSE) :  5114.916514601162
-----
R2 Score :  0.7434776049643312
-----
```

```
[186]: ytest_pred = lin_reg_model.predict(xtest)
ytest_pred
```

```
[186]: array([ 9336.84134219,  7305.55652027, 32829.43460237,  9728.26767656,
   23557.46581852,  9532.63066135,  1692.8228067,  15446.45811413,
   2848.20365207, 11087.73846789, 24998.88123828, 9288.50796245,
   5716.57831586, 33088.43272113, 34743.77112981, 32096.96890174,
  13684.57022394, 31260.74817834, 9269.87475494, 28006.92562062,
   4718.53566847, 9648.92588828, 2847.19105337, 6596.73755658,
  11272.02930014, 12323.04621164, 13417.12864002, 6509.36221589,
  10120.58320588, 2836.60959399, 9082.43811341, 12539.04220522,
   4686.93299424, 4818.95326244, 4758.31016028, 11794.10344382,
  2874.8468731, 8723.90687444, 29841.12623571, 28324.64864922,
  4495.23886739, 4703.63709011, 13052.62833026, 11464.66547311,
  8347.97068685, 11767.31874966, 5472.24321934, 4104.657937,
  30896.85391013, 8489.78023707, 14657.54468937, 3055.98117503,
  11214.46327831, 2478.96394981, 12539.83561298, 11979.64385661,
  5083.27089836, 28059.6482575, 12502.46622191, 12223.18504359,
  13354.68834466, 9421.53049794, 14974.10403291, 8326.3900963,
  11452.78295875, 4932.99635378, 23622.77725862, 10977.00229023,
  3465.37379845, 5608.08101822, 9954.01150139, 11188.76108703,
  10536.66228178, 8949.41457436, 11005.31464239, 7004.52704208,
  7200.69135185, 10903.42598838, 6658.04213524, 9016.2867278,
  4022.22064191, 32520.71834864, 6569.75323139, 26709.62951322,
  27862.38966508, 30983.09798169, 7101.99659179, 12211.62949932,
  9853.85454914, 13250.80171625, 16182.50516535, 31664.32945181,
  29425.40790604, 65999.79486296, 27760.947153, 9098.5641489,
  25978.4722504, 4031.51329112, 25470.20477046, 6651.06214334,
  5954.67209627, 2590.08883126, 10422.60693844, 13825.88760868,
  11791.23602536, 4383.40177317, 9721.94592364, 28577.95589271,
  1419.07666262, 29507.89892963, 3461.94844874, 8800.25223323,
  13571.4195753, 27759.49672384, 11136.91401409, 4063.51422384,
  12686.37108612, 28197.66454847, 8223.91826559, 4024.80555736,
  8254.65391421, 10110.46136884, 13950.9633753, 5516.84422954,
  4851.81987486, 9935.32679945, 10309.09624516, 10836.18082189,
  13488.51483335, 6664.71288617, 5860.54285774, 8876.35480759,
  8916.7325924, 11552.44746323, 8070.93743258, 13953.12282207,
  8000.12709128, 27845.57918396, 30983.62371105, 27794.1710621,
  6591.75487787, 11915.3109435, 5875.94811258, 13478.02951434,
  3442.28572139, 29240.63787671, 6737.56004129, 5391.34583425,
  13383.51327857, 6938.943686, 33758.6193611, 3887.58676128,
  5341.1767603, 27605.12684035, 10560.50353343, 8099.98825933,
  13945.9807094, 9800.06609975, 23634.58342677, 29247.59994391,
  13836.32874118, 2822.94608458, 12877.61940056, 3259.46437644,
  5473.52026647, 10740.35598207, 34505.71484633, 32195.29748445,
  29841.7177913, 4286.71163422, 8143.6504353, 8397.322338,
  11065.15431875, 5537.49654412, 2869.18042222, 28621.66000451,
  22040.10140684, 15754.49383321, 24193.71818388, 18570.68010842,
  33072.72517848, 1527.24017684, 7762.55546405, 7522.33142282,
  5367.16052475, 5306.89945012, 6200.05441939, 4842.93945551,
  13431.68849089, 11043.59044813, 6475.34713299, 3255.56490361,
  2811.4258163, 27543.31953854, 14919.52218857, 11400.91183796,
  2541.75149818, 12247.214665, 2606.65431961, 8957.66051157,
  3235.43474222, 30097.43728402, 9587.55755036, 3479.29334785,
  22831.3147095, 23568.25865914, 9614.1794375, 3284.63192287,
  12005.06426069, 2518.36763691, 10672.66098384, 10869.26103671,
  14601.80307094, 23780.14481502, 7500.99785778, 5730.08973536,
  6334.60339528, 12994.52867996, 11288.72796688, 8146.35623478,
  5158.72841558, 11221.30841482, 12115.71093188, 30978.03038905,
  4481.98862983, 25291.34959456, 700.27906032, 3446.97313049,
  10784.4836049, 14198.45040041, 5053.32541998, 8024.78134016,
  4696.84617097, 27574.30903555, 8319.3156834, 11483.33739565,
  5908.74974817, 8971.43192439, 31358.77587333, 4581.68989826,
  10434.72264793, 26901.88262948, 5796.10696856, 5102.78547631,
  2214.0430209, 5162.62788841, 5127.49786633, 6477.70175567,
  16484.26377745, 461.9844441, 2979.11366305, 9969.44471179,
  3951.04582982, 10632.63725819, 4207.28866377, 5399.10129639,
  11921.03384572, 6206.22975509, 8775.95237362, 7727.82639383,
  9105.37856848, 11013.38547821, 24485.65586146, 34105.2803857,
  11708.79820287, 7899.1279854, 34900.71316165, 11540.1112273,
  8098.55453728, 8812.07217133, 9526.71870311, 10668.74017653,
  10676.47774352, 15750.26159333, 2231.65423443, 21248.37476425,
  11365.03359427, 28883.12538138, 6099.0553135, 12157.7633065,
  10269.28887046, 16183.60375609, 9973.65170622, 10933.6309583,
  28031.24098256, 4108.09729653, 12622.11028552, 33896.37278921,
  5975.46029543, 5971.24153796, 3882.84513481, 11049.95365181,
  23367.02315102, 13110.06290588, 9683.01410007, 10270.25444265,
  12422.99390998, 2520.32299755, 3363.74463598, 26301.29317323,
  26190.58324457, 11929.84950227, 4320.78536005, 22611.07522761,
  12941.18415136, 26627.23467246, 4039.3067653, 34059.42128225,
  10308.26292815, 4774.41896126, 7175.22972777, 3206.25520501,
  23161.636598, 13443.80815371, 1856.89671964, 12244.28470488,
  11984.59859651, 13828.57277726, 30299.64226249, 13553.01889453,
  28231.20096171, 9668.76026547, 16275.83592217, 5995.64639356,
  9194.73773995, 9220.22955355, 14341.93454043, 9482.61806502,
  7631.75305592, 13458.72149538, 12532.86686952, 13654.68780122,
  7232.98262683, 22898.96629933, 8894.82096699])
```

```
[187]: # for Testing

print("----- Evaluation for Testing Data -----")
mse = mean_squared_error(ytest, ytest_pred)
print("mean squared error (MSE) : ", mse)
print("-"*80)

mae = mean_absolute_error(ytest, ytest_pred)
print("mean absolute error (MAE) : ", mae)
print("-"*80)

rmse = np.sqrt(mse)
print("Root mean squared error (RMSE) : ", rmse)
print("-"*80)

r2score = r2_score(ytest, ytest_pred)
print("R2 Score : ", r2score)
print("-"*80)
```

Evaluation for Testing Data

```
Evaluation for Testing Data
mean squared error (MSE) : 25663562.89341267
-----
mean absolute error (MAE) : 3400.935136082026
-----
Root mean squared error (RMSE) : 5065.9217219981465
-----
R2 Score : 0.7589437714409245
```

Algorithm 2: L1- Lasso

```
[194]: Lasso_model=Lasso()
Lasso_model

[194]: + Lasso ⓘ ?  
Lasso()
```

```
[196]: # Training the Lasso model
lasso_model = Lasso_model.fit(xtrain, ytrain)
ytrain_pred = lasso_model.predict(xtrain)

print("----- Evaluation of Lasso Training Data -----")
mse_train = mean_squared_error(ytrain, ytrain_pred)
print("Mean Squared Error (MSE):", mse_train)
print("-" * 80)

mae_train = mean_absolute_error(ytrain, ytrain_pred)
print("Mean Absolute Error (MAE):", mae_train)
print("-" * 80)

rmse_train = np.sqrt(mse_train)
print("Root Mean Squared Error (RMSE):", rmse_train)
print("-" * 80)

r2_train = r2_score(ytrain, ytrain_pred)
print("R2 Score:", r2_train)
print("-" * 80)

----- Evaluation of Lasso Training Data -----
Mean Squared Error (MSE): 26162383.23163086
-----
Mean Absolute Error (MAE): 3346.9011439298597
-----
Root Mean Squared Error (RMSE): 5114.917715040083
-----
R2 Score: 0.743477484555913
```

```
[198]: # Testing the Lasso model
ytest_pred = lasso_model.predict(xtest)

print("----- Evaluation of Lasso Testing Data -----")
mse_test = mean_squared_error(ytest, ytest_pred)
print("Mean Squared Error (MSE):", mse_test)
print("-" * 80)

mae_test = mean_absolute_error(ytest, ytest_pred)
print("Mean Absolute Error (MAE):", mae_test)
print("-" * 80)

rmse_test = np.sqrt(mse_test)
print("Root Mean Squared Error (RMSE):", rmse_test)
print("-" * 80)

r2_test = r2_score(ytest, ytest_pred)
print("R2 Score:", r2_test)
print("-" * 80)

----- Evaluation of Lasso Testing Data -----
Mean Squared Error (MSE): 25666205.392001774
-----
Mean Absolute Error (MAE): 3401.5554279432567
-----
Root Mean Squared Error (RMSE): 5066.18252651854
-----
R2 Score: 0.7589189506182465
```

Algorithm 3: L2 - Ridge

```
[201]: Ridge_model=Ridge()
Ridge_model

[201]: + Ridge ⓘ ?  
Ridge()
```

```
[203]: ridge_model=Ridge_model.fit(xtrain,ytrain)
ytrain_predict = Ridge_model.predict(xtrain)
ytrain_predict
```

```
[203]: array([ 3044.79386716, 5204.46305323, 12733.99715313, ...,
       11472.52680989, 31410.39392572, 11024.914144 ])
```

```
[205]: print("----- Evaluation of ridge Training Data -----")
mse = mean_squared_error(ytrain, ytrain_predict)
print("Mean Squared Error (MSE):", mse)
print("-" * 80)
```

```
mae = mean_absolute_error(ytrain, ytrain_pred)
print("Mean Absolute Error (MAE):", mae)
print("-" * 80)
rmse = np.sqrt(mse)
print("Root Mean Squared Error (RMSE):", rmse)
print("-" * 80)

r2score = r2_score(ytrain, ytrain_pred)
print("R2 Score:", r2score)
print("-" * 80)

----- Evaluation of ridge Training Data -----
Mean Squared Error (MSE): 26162383.23163086
-----
Mean Absolute Error (MAE): 3346.9011439298597
-----
Root Mean Squared Error (RMSE): 5114.917715040083
-----
R2 Score: 0.743477484555913
-----
```

```
[207]: ytest_pred = Ridge_model.predict(xtest)
ytest_pred

print("----- Ridge Evaluation for Testing Data -----")
mse = mean_squared_error(ytest, ytest_pred)
print("mean squared error (MSE) : ", mse)
print("-" * 80)

mae = mean_absolute_error(ytest, ytest_pred)
print("mean absolute error (MAE) : ", mae)
print("-" * 80)

rmse = np.sqrt(mse)
print("Root mean squared error (RMSE) : ", rmse)
print("-" * 80)

r2score = r2_score(ytest, ytest_pred)
print("R2 Score : ", r2score)
print("-" * 80)

----- Ridge Evaluation for Testing Data -----
mean squared error (MSE) :  25690567.64871225
-----
mean absolute error (MAE) :  3412.9725338121752
-----
Root mean squared error (RMSE) :  5068.586356047634
-----
R2 Score :  0.7586901174766364
-----
```

```
[ ]:
```

